

Sudoku 4x3 GPU Exact Enumeration Research Log

Verification

Verification of 2006 Pettersen/Silver result used 1035 CPU-hours over a twelve-day period between 8 June 2022 and 20 June 2022.

Note on 3x4 vs 4x3

I have been using 3x4 and 4x3 somewhat interchangeably. I starting using just 3x4 because that's how the Wikipedia table of results had it. After reconnecting with Kjell after all these years, I saw that he used 3x4 and 4x3 to refer to two distinct methods for the exact enumeration.

- In a 4x3 count, the 144,578 gangsters (equivalence class representatives) of a four-box, 12x4 band are determined and used with the other two bands.
- In a 3x4 count, the 2230 gangsters of a three-box, 12x3 stack are determined and used with the other three stacks.

Both the original 2006 enumeration, and this 2022 verification, are 4x3. Both methods must give the same result, but both Kjell and I believe that 4x3 is more efficient. Kjell has recently been thinking about the 3x4 version, and he may yet develop an efficient way to do it.

I think Kjell is right that I should be naming all this work 4x3, not 3x4, but there is a lot of 3x4 in filenames and the like because that's where I started. If you see 3x4, keep in mind that the method under investigation is 4x3.

Workload

CPU-hours is easy to tally but not a great way to measure the workload. It's very dependent on the machines I happened to have available, and tends to be biased by the slowest ones. It is a measure of serial hours, i.e. the time that would have been needed if the machines were not run in parallel. It also was affected by some difficulty I had in getting Windows and Ubuntu to run the cores at max speed, instead of trying to conserve power.

A parallel measure of CPU-hours would be easy to tally if all machines were run for the entire time—it would be essentially the same as the formula for parallel resistance. But that was not the case in practice.

Parallel execution on multiple cores of a single CPU is not the same as parallel execution on separate machines, because the parallel threads compete with each other for shared resources. Each of the 144,578 gangsters is enumerated independently, and the time for each one is defined and recorded as the elapsed time divided by the number of parallel threads.

Computers

The first six were used in the verification run. The speed is the average time to enumerate one gangster with a 32-gangster benchmark run of 865 – 896 (the first 32 in group 1).

Name	GHz	Cores	Threads	CPU	OS	Compiler	Speed
PT2017	3.10	4	8	x64 Xeon E3-1535M v6	Windows	MS VC++	16.5
PT2019	2.11	4	8	x64 i7-8650U	Windows	MS VC++	24
Judy7		4	8	x64	Windows	MS VC++	
PT2015		2	4	x64	Ubuntu	GCC	
Judy6		2	4	x64	Windows	MS VC++	
CGNX	2.90	2	4	x64 i7-7600U	Windows	MS VC++	
EPT2022	2.26	8	8	ARM-64 v8.2	Ubuntu	GCC	10

The speed (actually seconds/gangster) numbers are somewhat variable, run to run. For example, most of the PT2019 runs are in the 23.8 – 24.5 range, but a small number came in at around 22.1. I don't understand this 2 seconds/gangster bimodal variation on this particular machine. My current speculation is that the variation in parallel thread order may interact with the data caches and hyperthreading and occasionally produce this weird bimodal behavior.

More detailed timing follows. This version is slightly different than the baseline verification run. The order of the DoubleBoxCount outer loop was modified to try to achieve slightly better data cache performance. It may have made a very small improvement in speed, barely measurable.

PT2017

Using 8 threads

1,180,382 cache misses 4,087,416 code calls

Read count file gridCount_1-.txt, total time so far 0.46 hours

Profile tree:

```

Sudoku3x4                                     535.947
  RowCode Init                               15400 *          0.0661 -> 0.001
  ColCode Init                               5775 *           0.0579 -> 0.000
  Row Tables                                369600 *          0.1185 -> 0.044
  Column Tables                             138600 *          0.1487 -> 0.021
  BoxCompatible Init                          715 *           1.6134 -> 0.001
  Column Nodes                               31104 *          14.3970 -> 0.448
  BandGang Construct                          0.247
  Verify BandGang Tables                     0.019
  Band Gangsters                             3.790
  Fix gang cache                             300155625 *         0.0022 -> 0.672
  Read/verify gangsters                       144578 *          1.2478 -> 0.180
  Replace cache codes                         300155625 *         0.0021 -> 0.645
  Construct GridCounter                       1998150 *          0.2037 -> 0.407
  Grid counter                               528.615
    GridCounter Setup                         0.906
      Big tables                             119716 *          7.3130 -> 0.875
      Sort                                   0.026
      Overhead                               0.005
    Main count loop                           32 * 16490904.0063 -> 527.709
    Overhead                                 0.000
  Overhead                                   0.856

```

PT2019

Using 8 threads

1,175,430 cache misses 4,070,668 code calls

Read count file test1_1-.txt, total time so far 0.00 hours

Profile tree:

Sudoku3x4				772.867
RowCode Init	15400 *	0.0625 ->		0.001
ColCode Init	5775 *	0.0472 ->		0.000
Row Tables	369600 *	0.1273 ->		0.047
Column Tables	138600 *	0.1496 ->		0.021
BoxCompatible Init	715 *	1.6615 ->		0.001
Column Nodes	31104 *	15.4951 ->		0.482
BandGang Construct				0.256
Verify BandGang Tables				0.020
Band Gangsters				5.599
Fix gang cache	300155625 *	0.0023 ->		0.697
Read/verify gangsters	144578 *	1.2698 ->		0.184
Replace cache codes	300155625 *	0.0023 ->		0.676
Construct GridCounter	1998150 *	0.2092 ->		0.418
Grid counter				763.586
GridCounter Setup				0.951
Big tables	119716 *	7.6898 ->		0.921
Sort				0.030
Overhead				0.000
Main count loop	32 * 2383	2356.1656 ->		762.635
Overhead				0.000
Overhead				0.879

Using 8 threads

8.5697 [144578]; 1,176,526 cache misses 4,075,140 code calls

1,176,526 cache misses 4,075,140 code calls

Profile tree:

Sudoku3x4				719.754
RowCode Init	15400 *	0.0735 ->		0.001
ColCode Init	5775 *	0.0550 ->		0.000
Row Tables	369600 *	0.1282 ->		0.047
Column Tables	138600 *	0.1622 ->		0.022
BoxCompatible Init	715 *	1.8829 ->		0.001
Column Nodes	31104 *	14.8154 ->		0.461
BandGang Construct				0.259
Verify BandGang Tables				0.029
Band Gangsters				5.337
Fix gang cache	300155625 *	0.0023 ->		0.702
Read/verify gangsters	144578 *	1.2762 ->		0.185
Replace cache codes	300155625 *	0.0023 ->		0.676
Construct GridCounter	1998150 *	0.2091 ->		0.418
Grid counter				710.708
GridCounter Setup				0.963
Big tables	119716 *	7.8065 ->		0.935
Sort				0.028
Overhead				0.001
Main count loop	32 * 2217	9517.5406 ->		709.745
Overhead				0.000
Overhead				0.909

EPT2022

Using 8 threads

1,158,310 cache misses 4,009,004 code calls

Read count file gridCount_1-.txt, total time so far 0.46 hours

Profile tree:

Sudoku3x4				326.179
RowCode Init	15400 *	0.1489 ->		0.002
ColCode Init	5775 *	0.0968 ->		0.001
Row Tables	369600 *	0.2255 ->		0.083
Column Tables	138600 *	0.3447 ->		0.048
BoxCompatible Init	715 *	4.3080 ->		0.003
Column Nodes	31104 *	23.6487 ->		0.736
BandGang Construct				0.250
Verify BandGang Tables				0.022
Band Gangsters				4.772
Fix gang cache	300155625 *	0.0019 ->		0.581
Read/verify gangsters	144578 *	0.5562 ->		0.080
Replace cache codes	300155625 *	0.0022 ->		0.649
Construct GridCounter	1998150 *	0.2685 ->		0.536
Grid counter				317.587
GridCounter Setup				2.208
Big tables	119716 *	6.0924 ->		0.729
Sort				1.478
Overhead				0.001
Main count loop	32 * 985	5592.8388 ->		315.379
Overhead				0.000
Overhead				0.829

Jetson AGX Xavier GPU Basics

Clock rate 1377000 kHz

L2 cache size 524288

Max blocks per multiprocessor 32

Max grid size 2147483647.65535.65535

Max block dimension 1024.1024.64

Max threads per block 1024

Max threads per multiprocessor 2048

Multiprocessor count 8

Reserved shared memory per block 0 bytes

Shared memory per block 49152 bytes

Shared memory per multiprocessor 98304 bytes

Total global memory on device 32517738496 bytes

Warp size in threads 32

Cuda First Cut

First cut at a Cuda GPU program for executing the main grid counting loop for the Sudoku 4x3 exact count. It makes very poor use of GPU resources, and is actually slower than running the CPU code on the Jetson's 8-core ARM v8.2 processors. The purpose of this first cut is to confirm that I understand the Nvidia tool chain and the most basic operations of Cuda. The code runs and gets the correct results.

The first step in using the GPU properly will be to deal with the poor memory access pattern, which radically degrades the GPU's memory bandwidth and stalls the compute elements.

Here is a pure GPU run—one thread runs all the setup, and calls Cuda code to run 16 blocks of 32 threads to do the main counting loops.

Using 1 thread
 1,052,036 cache misses 3,653,317 code calls
 Read count file gridCount_1-.txt, total time so far 0.46 hours

Profile tree:

Sudoku3x4			389.234
RowCode Init	15400 *	0.1559 ->	0.002
ColCode Init	5775 *	0.1008 ->	0.001
Row Tables	369600 *	0.2489 ->	0.092
Column Tables	138600 *	0.3945 ->	0.055
BoxCompatible Init	715 *	5.3005 ->	0.004
Column Nodes	31104 *	23.1097 ->	0.719
BandGang Construct			0.252
Verify BandGang Tables			0.028
Band Gangsters			10.687
Fix gang cache	300155625 *	0.0019 ->	0.580
Read/verify gangsters	144578 *	0.5620 ->	0.081
Replace cache codes	300155625 *	0.0019 ->	0.562
Construct GridCounter	1998150 *	0.2820 ->	0.564
Give band counts to GPU	300155625 *	0.0026 ->	0.795
Grid counter			373.769
GridCounter Setup			2.582
Big tables	119716 *	6.3462 ->	0.760
Sort			1.616
Tables -> GPU			0.205
Overhead			0.001
Main count loop	32 * 1159	9569.5722 ->	371.186
Overhead			0.000
Overhead			1.045

Here is a heterogeneous run—seven threads run on the ARM cores and one thread calls the

Cuda/GPU code:

Using 8 threads
 1,158,288 cache misses 4,009,055 code calls
 Read count file gridCount_1-.txt, total time so far 0.46 hours

Profile tree:

Sudoku3x4			336.808
RowCode Init	15400 *	0.1379 ->	0.002
ColCode Init	5775 *	0.0980 ->	0.001
Row Tables	369600 *	0.1976 ->	0.073
Column Tables	138600 *	0.2959 ->	0.041
BoxCompatible Init	715 *	3.6988 ->	0.003
Column Nodes	31104 *	24.4196 ->	0.760
BandGang Construct			0.253
Verify BandGang Tables			0.022
Band Gangsters			4.569
Fix gang cache	300155625 *	0.0020 ->	0.589
Read/verify gangsters	144578 *	0.5550 ->	0.080
Replace cache codes	300155625 *	0.0021 ->	0.629
Construct GridCounter	1998150 *	0.2684 ->	0.536
Give band counts to GPU	300155625 *	0.0028 ->	0.854
Grid counter			327.460
GridCounter Setup			36.896
Big tables	119716 *	6.1662 ->	0.738
Sort			35.850
Tables -> GPU			0.306
Overhead			0.001
Main count loop	32 * 908	0136.8877 ->	290.564
Overhead			0.000
Overhead			0.935

Compilation Command Line

```
nvcc --m64 --std c++17 --compiler-options -std=c++17,-march=armv8-a+simd,-Ofast,-wno-format,-DJETSON --linker-options -pthread --include-path . -o sudoku3x4 bignumMT.cpp profile.cpp general.cpp timer.cpp Sudoku3x4.cpp gridCount.cu
```